

Name 1:

Name 2:

COMPUTER NETWORKING

LAB EXERCISES (TP) 2

L1 v.s. L2 v.s. L3, NAT AND TROUBLESHOOTING

With Solutions

October 14, 2016

Abstract

In this Lab you will work with the virtual environment introduced in Lab 1. First you will see the different behavior of networking devices that work on Layer 1, layer 2 and layer 3; then you will configure your virtual network to be able to access the Internet; later you will help Jon and Arya to fix their networking problems when a common enemy, Joffrey, changes the configuration in their network. Finally you will practice reading TCP-IP headers to collect and correlate information.

1 PREPARING THE LAB

1.1 LAB REPORT

Type your answers in this document. We recommend you use Adobe Reader XI to open this PDF. When you finish, save the report and upload it on moodle. Don't forget to write your names on the first page of the report. **The deadline is Wednesday, October 26th, 23:59:59**

1.2 SETTING UP VIRTUAL MACHINE

In this Lab, you will work with the same virtual machine that you created in Lab 1. Copy the **lab2** folder from Moodle into the shared folder of your VM before starting the lab.

1.3 USING SCRIPTS

As a general advice, use scripts to save your work for each section, especially Section 3. This is useful for 1) saving time and not repeating the same commands each time you restart Mininet, and 2) reviewing and debugging your work in case you run into issues.

2 LAYER 1 VS. LAYER 2 VS. LAYER 3 NETWORKING

The aim of this section is to illustrate the difference between networking devices that work at layer 1, layer 2 and layer 3. For this exercise we only consider IPv4 addressing.




Figure 1: Network configuration with two hubs in the center

2.1 USING HUB AS A NETWORKING DEVICE

A hub is layer-one intermediate system that repeats bits (symbols) received from one of its ports to all other ports. In this section we analyze how it works.

Open a terminal in your VM and run the script `topol.py`, which should be located in the shared folder on the Desktop. If not, refer to Section 1.2.

```
# python topol.py
```

This will create the network described in Figure 1, and redirect you to the mininet CLI. Additionally, one terminal will appear for each of the four hosts. The four new terminals will be labeled (h1, h2, h3, h4) for convenience.

Run the following two commands to configure s12 and s34 as hubs.

```
mininet> sh ovs-ofctl add-flow s12 action=flood
mininet> sh ovs-ofctl add-flow s34 action=flood
```

h1, h2, and h3 should be located on the 10.0.0.0/24 subnet with the fourth byte of their IP address being 1, 2, and 3, respectively. h4 should have the IP address 10.0.1.4/24. Check the configuration of the IP addresses for each of the hosts in their respective terminals and correct any wrong configuration.

Q1/ Was there a wrongly configured host? How did you find out?

Solution. *h3 was configured with a wrong subnet mask. ip addr show would help us find out.*

Q2/ Which line of code in the mininet configuration script does this error correspond to?

Solution. Line 23 in the `topo1.py` file should be:

`h3 = net.addHost('h3', ip='10.0.0.3/24')`

instead of

`h3 = net.addHost('h3', ip='10.0.0.3/31')`

Q3/ What command did you run to fix the configuration issue?

Solution. In the `h3` terminal: `ip addr flush dev h3-eth0; ip addr add 10.0.0.3/24 dev h3-eth0`

Start Wireshark on all four hosts. It will be hard to keep track of which Wireshark window corresponds to which host. One way to do so would be to start Wireshark on the hosts in order, i.e. `h1`, then `h2`, then `h3`, and finally on `h4`. This way the Wireshark windows will be in this same order in the taskbar. Start capturing on all the `eth0` interfaces.

```
# wireshark &
```

From `h1`, ping `h2`.

```
mininet> h1 ping h2
```

Q4/ What would be another way of doing the same thing?

Solution. From the `h1` terminal, ping `10.0.0.2`

Q5/ Is there any difference between the traffic captured by the four hosts in Wireshark?

Solution. No, In all machines we can see exactly the same packets

Q6/ Explain why you see these results.

Solution. In a hub, the incoming traffic from a port is forwarded to all ports in the hub (except the one it received the packet from) without any type of filtering. A hub is an intermediate system that just amplifies and repeats signals.

2.2 USING A BRIDGE AS A NETWORKING DEVICE

A bridge is a link-layer intermediate system which expands a LAN by making forwarding decision based on destination MAC-address. In this section you will learn how they work.

We will first change `s34` to act as a bridge by running the following command.

```
mininet> sh ovs-ofctl add-flow s34 action=normal
```

Note that `s12` and `s34` were originally created as bridges, but we modified them in the previous section to act as hubs by adding flows. You will learn more about flows in Lab 4.

Q7/ If you were to exit Mininet and run `topo1.py` again, what command(s) would you then execute to reach the topology we have now?

Solution. From the mininet CLI; `ovs-ofctl add-flow s12 action=flood`

Now, let's test our bridge configuration. Start a Wireshark capture on all four hosts again, and again ping h2 from h1.

Q8/ Describe the different types of packets observed on h1, h2, h3 and h4.

Solution. On h1 and h3 we are able to see ICMP echo-request, ICMP echo-reply, ARP requests and ARP replies. In h2 we only observed ARP request packets (broadcast)

Q9/ Explain the results. What is the difference compared to having two hubs?

Solution. In a bridge, frames are forwarded by searching in the MAC-address table and performing an exact match lookup. If the search finds a match, the frame is forwarded directly to the port given by the MAC-address table. If the frame is not in the table, then the bridge forwards the frame to all ports except the one it received the frame from.

This is why we only receive broadcast packets on h2 (destination MAC-address of broadcasts is not known), but the actual ICMP traffic we do not see it (the bridge sends traffic directly from h1 to h3 and vice versa)

Q10/ Ping from h1 to h4. Observe the traffic captured and explain your findings.

Solution. We don't see any packets. h1 will use its network mask on h4's IP address and check if they are in the same subnet. As they are not in the same subnet, h1 will attempt to contact its default gateway to send the packet, and if no default gateway is configured (which is this case), it will not send any packet at all.

Now, focus on the interfaces of h1, h2 and h3. Ping from h1 to h3.

Q11/ Observe the packets captured on h2 and explain the results.

Solution. We observe the ICMP requests and replies on h2, since h1 and h2 are still connected by a hub.

Q12/ Compare the packets sent by h1 to the ones received by h3, specifically at source/destination MAC-addresses. Explain the similarities and differences, if any.

Solution. Traffic is the same, same ethernet header, same source/destination MAC-addresses. The Ethernet bridge does not affect any source/destination MAC-address, it is transparent to the MAC and IP layers.

2.3 USING A ROUTER AS A NETWORKING DEVICE




Figure 2: Network configuration with a router and a switch

We have already configured a router in Lab 1, but we did not address how it worked. In this section we learn about the process of routing a packet.

We'll begin by exiting Mininet, cleaning up the previous topology, and running `topo2.py`. The new topology consists of three hosts, a router, and a bridge, as shown in Figure 2.

```
mininet> exit
# mn -c
# python topo2.py
```

Perform a reachability test in Mininet. A reachability test is a test to determine which hosts can 'reach' one another. This is performed by having each host ping all other hosts. In our case, this also includes the router. A quick way to do this test in Mininet is by running the following command.

```
mininet> pingall
```

Q13/ What is the percentage of dropped packets? What does this correspond to?

Solution. *33%, one third of the links are broken*

Q14/ Which hosts are unable to reach one another?

Solution. *h1 and h3; h2 and h3*

We will now attempt to fix the problem. First, open the `topo2.py` script and inspect it.

Q15/ What is the subnet mask used throughout the file?

Solution. *255.255.255.0*

Q16/ What are the interfaces and respective IP addresses of the router `r1`?

Solution. *r1-eth0: 10.0.0.100; r1-eth1: 10.0.1.100*

Q17/ Can you spot any misconfigurations in the file?

Solution. *disabled ip forwarding; default gateway of h2 is wrong*

The first issue we notice in the file is that IP forwarding has been disabled in the router `r1`. We will enable it and attempt the reachability test again.

Q18/ What is the command you will use to enable IPv4 forwarding on `r1`?

Solution.

```
# echo 1 > /proc/sys/net/ipv4/ip_forward
```

Q19/ What is the percentage of dropped packets now in the reachability test? Which hosts are still unable to ping each other?

Solution. *16%; h2 and h3*

The next issue in the file is the default gateway of h2, which is set to 10.0.0.101. This IP address does not exist in our network. Delete this configuration and set a proper default gateway through the terminal of h2.

Q20/ What are the commands you used to achieve this?

Solution. *ip route del default via 10.0.0.101; ip route add default via 10.0.0.100*

Q21/ What are the results of the reachability test now?

Solution. *0% dropped packets. Our network is now properly configured.*

Now we learn about routing between two interfaces. Monitor the traffic in Wireshark of both interfaces of r1. From h1, ping h3.

Q22/ What changes are done to IP packets when they are routed between r1's r1-eth0 and r1-eth1?

Solution. *We see that the original source MAC-address of h1 is replaced with r1's r1-eth0 MAC-address and destination MAC-address of r1's r1-eth0 is replaced with h3's eth0. Also we see that TTL is decremented by 1.*

Q23/ What is the purpose of such changes?

Solution. *The reason is to adjust the IP packet to the new LAN where it is being routed to. When routing from eth0 to eth1 (or vice versa), r1 removes the MAC header, next it reads the IP header, then it makes forwarding decisions based on the destination IP address, and finally it inserts a new MAC header which has the source/destination MAC-addresses compatible with the scope of the LAN between r1 and h3. The reason for TTL is to avoid loops in the network.*

2.4 ROUTING WITH MULTIPLE HOPS

In this section we want to see what happens when we introduce a second intermediate routing device to the network. Let's start by setting the default gateway of r1 to h3's IP address, then on h3 remove the default gateway and enable IPv4 forwarding.

Q24/ Type in the commands you need to do this.

Solution. *For h3*

```
# ip route del default via 10.10.24.4
# echo 1 > /proc/sys/net/IPv4/ip_forward
```

For r1

```
#ip route add default via 10.0.1.3
```

Monitor eth0 of h3 and r1. Try pinging from h1 to h3 and h2 to h3.

Q25/ Based on Wireshark captures, explain why it does not work.

Solution. *From r1 we see ICMP echo-requests sent to h3. On h3 we don't see any ICMP echo-reply. h3 does not know how to reach h1's IP address, which means that there is no route configured in h3 to return ip packets to h1 and h2. By default, if a network device does not know how to route a packet, it drops it.*

Q26/ Write down **one single command** that fixes the problem (namely pinging from h1 and h2 to h3), and specify the PC where you need to apply it.

Solution. *One wrong solution would be to return the default route on h3, and point to r1, but this would disable any future Internet access on h3, thus even though it would solve the problem for now, it will not work in the future section.*

A second approach is to add static routes in h3 in order to reach h1 and h2:

```
#ip route add 10.0.0.0/24 via 10.0.1.100
```

Ping again from h1 and h2 to h3, and confirm that your fix solves the problem before moving to the next section.

3 CONNECTING VIRTUAL ENVIRONMENT TO THE REAL WORLD USING NETWORK ADDRESS TRANSLATION (NAT)

In this section we will use what we learned from Lab1 about manipulating the `iptables` filter. The purpose of the section is to connect the isolated virtual network that we have deployed so far, to the real Internet.

We will work in the network described in Figure 3. h1 and h2 are workstations, r1 is an aggregation router, and h3 is the perimeter router where we will have our connection to the real world.




Figure 3: Network configuration with a connection to the real world

Q27/ We have one real connection (IP address) to the real world, but we have two clients (h1 and h2) that require access to the internet. Which solution would you use to tackle this problem, and explain how would it solve the problem.

Solution. *We use network-address translations (NAT). NAT would create separate mappings by using separate TCP/UDP port of the real IP address of h3's h3-eth1 for each of h1's and h2's connections.*

There are two main steps to connect your virtual environment to the real Internet:

1. We require a real IP address on h3-eth1 interface of h3.
2. We need to masquerade the traffic coming from h1 and h2.

3.1 BORROWING AN IPV4 ADDRESS

As we said before, we need an *existing* (real) IPv4 address that can be used to connect to the Internet. The purpose of this section is to obtain such valid IPv4 address. There are three steps for that:

1. Create a bridge between a real interface in your host machine, and h3's eth1.
2. Finding a suitable IPv4 address.
3. Setting up the IPv4 address of h3-eth1 on h3.

3.1.1 BRIDGE BETWEEN THE PHYSICAL AND THE VIRTUAL INTERFACE

For the first step, the process is shown schematically in Figure 4, and we will cover it step by step.

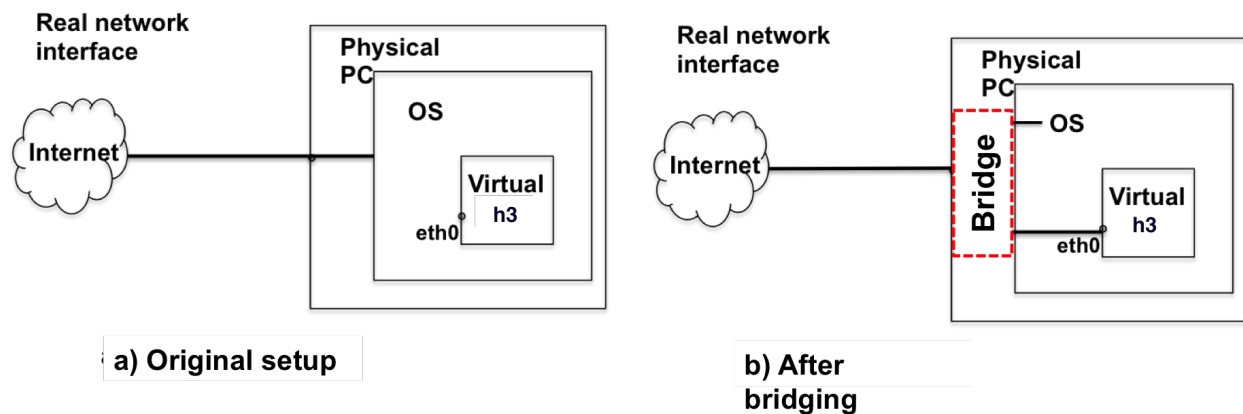


Figure 4: Bridging the network adapter

To perform the bridging between physical and virtual network-adapters, we first need to find out which interface our VM is using to access the internet.

Exit Mininet, clean up the topology, and run wireshark on the VM. Start capturing on each interface and ping `www.google.com` to find out if that interface is being used to access the internet.

```
mininet> exit
# mn -c
# wireshark &
# ping www.google.com
```

Q28/ What is the IP address of the interface used to access the internet? Do you notice anything special about it?

Solution. *Depends on the person. It is a private address given by the NAT of the VM, and not the public address given by the bridged adapter. The interface itself is important for the next few questions.*

Now that we know which interface is used by the VM, the next step is to implement a bridge between it and h3-eth1.

In Mininet, we can only implement a bridge with an OVS. To mitigate this, we will add an OVS between h3 and the internet cloud in Figure 3.

Run the script `topo2_internet.py`. This will create the same topology as before (in addition to all the extra configurations that you have performed, and add an extra OVS that we will use to perform the bridge.

From Mininet, execute the following command.

```
mininet> sh ovs-vsctl add-port s3 ethX
```

Replace X with the number of the interface that accesses the internet in your VM. Remember that the sudo password for the VMs is lca2. Now head over to the terminal of h3 and run the following command.

```
# dhclient h3-eth1
```

This automatically sets a usable IP address to the eth1 interface of h3, allowing it to access the internet through the bridge we just set up. Test the configuration by pinging Google or EPFL from h3.

3.1.2 CHOOSING THE BEST IP ADDRESS FOR h3'S ETH1

The task in this section is to find a suitable IP address for h3's eth1. This was basically achieved using the dhclient command in the previous section. Here, we will learn to do choose an appropriate IP address manually.

Run ifconfig and route -n get default (in Linux or Mac) or ipconfig /all (in Windows) and check your physical IP address, subnet mask and default gateway.

Q29/ Write down your physical IP address, subnet mask and default gateway. According to your network configuration, what is the range of IP addresses that we could use for the virtual adapter?

Solution. *Example output:*

IP Address: 128.178.151.219

Network Mask: 255.255.255.0

Default Gateway: 128.178.151.1

So, in principle we can choose anything in the 128.178.151.0/24 network except 128.178.151.1 and 128.178.151.219 assuming there is no other device using an IP address on that segment. Note that in this case 128.178.151.0 is the network address and 128.178.151.255 is the broadcast address, and therefore cannot be used as valid IP addresses.

Q30/ How could you find the current non-used IP addresses in your LAN?. Is it safe to take any of them?

Solution. *One solution would be by pinging the broadcast IP address, and then check in the ARP table all rows that have a complete IP - MAC address entry. Nonetheless, most network managers would not recommend this approach as it has high consumption of network resources. The only safe IP address you can take is your physical-adapter's IP address, any other IP address is susceptible of being allocated at any given moment by the DHCP service running on the EPFL network or on the network where you are doing the lab.*

Q31/ If you are doing this exercise outside EPFL (e.g. at your home), you will most likely get a private IP address and default gateway (e.g. 192.168.x.x or 10.x.x.x). Can you use these private IP addresses on h3's eth1 to do NAT? Explain why.

Solution. *Yes we can use it, it is called double NAT. The private IP address would behave as a "public" IP address from the Mininet point of view, and NAT would work as expected. In the outside world (from Mininet's perspective) there is another NAT box (e.g. your home ADSL router) that translates the private IP address you received into a valid public IP address in order to reach the internet. This is the reason why the industry is not in a rush towards IPv6, and they prefer to do double, triple or quadruple NATs.*

3.2 NAT CONFIGURATION

We worked with the command used to configure NAT in Lab1, `iptables -t nat`, which manages the table that contains rules regarding address translations. In this section, we will analyze how NAT works for different types of packets.

First let's see what happens when h1/h2 access to the Internet with their native IP address. Monitor with Wireshark the interface `eth1` of h3. From h1/h2, ping to Google and its IP address.

```
# ping -c 5 www.google.com
```

```
# ping -c 5 172.217.18.100
```

Q32/ Analyze the packets coming from h1/h2 and explain why you are unable to reach Google.

Solution. In interface `eth0` of h3 we see packets with a private IP address in the source field. Most likely the first hop router in the ISP network will drop the packet as it is sourced by a private IP address. In the case you are doing the exercise from home and you have a private IP address in `eth0` of h3, it is likely that the private IP address you use for Mininet environment is already allocated by your ISP, therefore your ISP will route it according to its own rules which do not cover your virtual environment.

Q33/ Propose the `iptables -t nat` command you need to properly configure NAT in h3.

Solution.

```
# iptables -t nat -A POSTROUTING -o h3-eth1 -j MASQUERADE
```

Test from h1 and h2 and you have Internet connectivity by pinging Google and test that you have successfully configured your router to do NAT.

Next, let's explore how NAT works!!.

Do `traceroute` to Google from h2 and then from h3, while capturing `eth0` and `eth1` traffic on h3 using Wireshark. Explore the difference in the traffic on both cases.

Q34/ When doing `traceroute` from h2, what is the difference in the packets captured on h3's `eth0` and `eth1`?

Solution. The source IP address is modified according to the NAT rule, replacing the IP address of h2 with the IP address configured on `eth1` in h3.

Q35/ Focus on the `traceroute` from h3. What is the difference in the packets as compared to h2?

Solution. Source and destination ports are different. Source and destination IP addresses are the same.

Q36/ Which field in the UDP packet is used to identify the (local) source IP address of h2 in order to properly forward incoming ICMP replies back to it?

Solution. This is done via UDP port, and NAT device keeps a track on the private source IP/port and the correspondent public IP/port.

Do `ping` to Google from `h1` and `h3`, while capturing the traffic on `h3` (both on `eth0` and `eth1`) using Wireshark. Explore the difference in the traffic in both cases.

Q37/ What is the difference in the request ICMP packets captured between packets sent from `h1` and packets sent from `h3` when capturing on the exit interface of each?

Solution. *The ICMP query ID is different*

Q38/ Conclude how the incoming ICMP replies are forwarded back to `h1` when doing `ping` from `h1`. In particular, which field in the request/reply ICMP packets was used to identify the (local) source IP address?

Solution. *ICMP query replies are forwarded back to `h1` based on the QueryID taken from the ICMP packet header. If we issue the `iptables -t nat -L -v -n` command, we will see this IP address to Query ID mapping. This is handled according to **RFC 5508***

4 TROUBLESHOOTING

As the title for this section suggests, in networking things will not always work out as expected. Before starting to work in this section, you will have to execute a script. This script will put a number of PCs into a problematic situation where something doesn't work. Your task is to find out what the problem is and propose a solution.

You should not perform any debug command or Wireshark capture in `h3`. Assume `h3` is a router controlled by your Internet service provider (ISP) and you don't have access to it..

4.1 ABOUT GRADING THIS SECTION

The points given for your answer mostly depend on your explanations about how you located the problem!, so describe precisely your steps to locate and diagnose the fault. You should use the scientific method when answering the problem. The methodology you should use is the following:

1. Pose a hypothesis
2. Run experiments to validate the hypothesis
3. If validation is OK exit, else loop (go back to 1. by posing another hypothesis)

In your answer you should write down all steps. Specially, you should also write down all hypotheses that later proved to be wrong. We want to see the path you took to reach your final conclusion!

More specifically for this Lab: What were the commands you executed to get there? Up to which point did the system work as expected? What were the actions that never got executed but were expected? What was the packet that did not reach its destination? Where and why did it get dropped/lost?

4.2 WE WERE HACKED!!!

Jon and Arya are roommates and close friends. Both of them are connected to the Internet through the same home router thus the same ISP. The configuration is the same as described in Figure 5, where `h1` is Jon's computer and `h2` is Arya's computer, `r1` is the home router, shared between Jon and Arya, and `h3` is the ISP's router.

Jon and Arya are foreign students at EPFL and they use Facebook a lot to communicate with their relatives in their home town, Winterfell. None of them is an expert in computers and networking. They have a




Figure 5: Troubleshooting configuration

common enemy, Joffrey, who plays jokes and is a Computer Science student and network expert. Joffrey told Jon and Arya that he hacked their computers (h1 and h2) as well their home router (r1) and that they would never navigate through Facebook again.

To simulate Joffrey's malicious attack, we prepared a script for h1, h2 and r1. You should download the scripts from the course web page on moodle. Unzip and extract the corresponding file:

```
# tar -xzvf lab2-scripts.tar.gz
```

We recommend you to extract all scripts to your shared folder between host and virtual machines.

Let's start with Jon and h1. Check that on h1 you have a h1.hack directory with the following content:

```
# ls lab2-scripts
h1.hack
```

Navigate to the lab2-scripts folder and execute the script for h1 on h1. Make sure you **run the script as a superuser in Linux**.

```
# cd lab2-scripts
# sudo ./h1.hack
```

Note that if you stop the simulation in Mininet, then you need to redo the previous section and make sure that h1 and h2 have a proper internet connection, then reload the script on h1.

Also, if you do not manage to solve one of the questions, you may want to restart Mininet, redo the previous section and make sure that all hosts have an internet connection, and then move on to the next question. This is where using scripts to save the commands you performed to achieve a working configuration would help a great deal. Ideally, solving a question means undoing the changes brought about by the hack, so you can start solving any of the questions in this section from the stable working configuration of Section 3.

Now, your mission is to help Jon find out why he cannot navigate to Facebook anymore. You are only allowed to have Wireshark captures and use debug commands in h1.

Open firefox on h1 by typing `firefox &` in the corresponding terminal.

Q39/ Are there any problems in Jon's PC (h1)? Enumerate them (if any) and write down how you find them and how you fix them.

Solution. *Problem is that h1 has a wrongly configured IP address, so we need to change it back to the correct one, and add the default gateway that is removed as a result of flushing the ip address. This can be discovered by observing that no packets are captured on h1-eth0 and h3*

After you fix the issue, confirm that Jon can now safely log into Facebook to update his status to "I know nothing about TCP/IP".

Now repeat the same procedure for Arya: execute the `h2_hack` script on the h2 terminal. Again you can use Wireshark on h2 to capture packets.

Q40/ Open Firefox on h2 and try to navigate to Facebook. Are there any problems in Arya's PC (h2)? Enumerate them (if any) and write down how you find them and how you fix them.

Solution. *We notice that Arya is sending DNS queries but not receiving replies, which means that the nameserver is wrongly configured. Fixing that solves the problem.*

While you were solving these issues on h1 and h2, Joffrey had enough time to hack r1. To simulate this, execute the `r1_hack` script on r1. Doing this, Joffrey managed to stop both Jon and Arya from accessing Facebook again. Your job is to get both of them Facebook access again. Hurry though, Winter is Coming!

For this question, you may perform debugging and capture Wireshark scripts on h1, h2, and r1. DO NOT TOUCH h3!

Q41/ Enumerate the problems (if any) that were stopping Jon (h1) from accessing Facebook. How did you find them and fix them?

Solution. *All packets sent from from h1 to facebook IP addresses through r1 have a gateway assigned to them which routes them back where they came from. This can be observed by observing Wireshark capture on r1. Fixing this allows h1 to access the internet.*

However, when Jon tries to access Facebook, he is instead rerouted to the EPFL website (may not show the webpage on firefox due to security certificates). The problem is some iptables rules on r1 that change the destination of packets from Jon accessing Facebook to EPFL. They can be seen by running `iptables -t nat -v -L -n`, and deleted by running `iptables -t nat -F`.

Q42/ Enumerate the problems (if any) that were stopping Arya (h2) from accessing Facebook. How did you find them and fix them?

Solution. *The only issue with Arya is the routing table rules in r1 that was also an issue with Jon. After clearing this issue by removing these entries in the routing table (one by one), Arya will be able to access the internet (and Facebook), but Jon still had to fix the iptables issue as explained above.*

5 READING IPV4 AND IPV6 PACKET HEADERS

In this section, you will observe IPv4 and IPv6 connections in different contexts. Your task is to answer the questions based only on the traces presented for each connection. You no longer need Mininet for the remainder of this lab. However, you will use your VM for some of the questions.

5.1 IPV4 PACKET HEADERS

Jon is sitting in front of `lrcpc3` workstation and connects to `smartgrid.epfl.ch`. An ethical hacker has read all the frames passing on the network. Here are two packets resulting from this activity:

```
ETHER: ----- Ether Header -----
ETHER:
ETHER: Packet 2 arrived at 13:45:24.14
ETHER: Packet size = 74 bytes
ETHER: Destination = 0:0:d:2:ff:e2
ETHER: Source       = 0:0:f3:a4:43:ee
ETHER: Ethertype = 0800
ETHER:
IP: ----- IP Header -----
IP:
IP: Version = 4
IP: Header length = 20 bytes
IP: Type of service = 0x00
IP:     xxx. .... = 0 (precedence)
IP:     ...0 .... = normal delay
IP:     .... 0... = normal throughput
IP:     .... .0.. = normal reliability
IP: Total length = 60 bytes
IP: Identification = 2947
IP: Flags = 0x0
IP:     .0.. .... = may fragment
IP:     ..0. .... = last fragment
IP: Fragment offset = 0 bytes
IP: Time to live = 64 seconds/hops
IP: Protocol = 17
IP: Header checksum = c2ba
IP: Source address = 128.178.156.17
IP: Destination address = 128.178.25.8, IP:   No options
IP:
UDP: ----- UDP Header -----
UDP:
UDP: Source port = 1304
UDP: Destination port = 53 (DNS)
UDP: Length = 40
UDP: Checksum = B281
UDP:
DNS: ----- DNS: -----
DNS:
DNS: ""
DNS:
```

```
ETHER: ----- Ether Header -----
ETHER:
ETHER: Packet 3 arrived at 13:45:24.85
ETHER: Packet size = 202 bytes
ETHER: Destination = 0:0:c0:b8:c2:8d, Western Digital
ETHER: Source       = 0:0:c:2:78:36, Cisco
ETHER: Ethertype = 0800
ETHER:
IP: ----- IP Header -----
IP:
IP: Version = 4
IP: Header length = 20 bytes
IP: Type of service = 0x00
```

```

IP:      xxx. .... = 0 (precedence)
IP:      ...0 .... = normal delay
IP:      .... 0... = normal throughput
IP:      .... .0.. = normal reliability
IP:      Total length = 188 bytes
IP:      Identification = 38579
IP:      Flags = 0x0
IP:      .0.. .... = may fragment
IP:      ..0. .... = last fragment
IP:      Fragment offset = 0 bytes
IP:      Time to live = 58 seconds/hops
IP:      Protocol = 17
IP:      Header checksum = 3d0a
IP:      Source address = 128.178.25.8,
IP:      Destination address = 128.178.156.17,
IP:      No options
IP:
UDP:      ----- UDP Header -----
UDP:
UDP:      Source port = 53
UDP:      Destination port = 1304
UDP:      Length = 168
UDP:      Checksum = 0000
UDP:
DNS:      ----- DNS: -----
DNS:
DNS:      " "
DNS:

```

Q43/ What is the purpose for this packet exchange?

Solution. *Packet 1: lrcpc3 issues an UDP packet and searches for the IP address of smartgrid.epfl.ch. This packet is sent to the DNS server.*

Packet 2: DNS server answers with the requested IP address to lrcpc3.

Q44/ Within each packet, what accounts for the difference in the *length* field in different headers?

Solution. *The length of the headers is subtracted after each packet is decapsulated at each layer. 14 bytes for the Ethernet header and 20 bytes for the IP header*

Q45/ How many routers would you say are there between Jon's PC and the DNS server?

Solution. *Assuming that DNS server issues the packet with the same TTL as lrcpc3 (that is 64), we can conclude that the routers between the DNS and lrcpc3 have reduced the TTL by 6. This means that it passed through 6 routers in this exchange, but there is no guarantee that it will follow the same path in future exchanges.*

Q46/ Repeat the packet exchange on your VM and observe the packets captured on Wireshark. What fields do you expect to be different?

Solution. *In query packet: Source MAC address; Source IP address; IP header checksum; udp source port; udp checksum*

In reply packet: destination MAC address; destination IP address; IP header checksum; up destination port; udp checksum

5.2 IPV6 PACKET HEADERS

Arya is sitting in front of lrcpc3 workstation and visits the website www.ethz.ch. An ethical hacker has read all the frames passing on the network. Here are the first two packets resulting from this activity:

```
ETHER: ----- Ether Header -----
ETHER:
ETHER: Packet 1 arrived at 11:55:22.298
ETHER: Packet size = 86 bytes
ETHER: Destination = 33:33:ff:01:00:01
ETHER: Source = 3c:07:54:3e:ab:f2
ETHER: Ethertype = 0x86dd
ETHER:
IP: ----- IP Header -----
IP:
IP: Version = 6
IP: Traffic class = 0x00000000
IP: .... 0000 00.. .... .... .... .... = Default Differentiated Service Field
IP: .... .... ..0. .... .... .... .... = No ECN-Capable Transport (ECT)
IP: .... .... ...0 .... .... .... .... = No ECN-CE
IP: .... .... 0000 0000 0000 0000 0000 = Flowlabel: 0x00000000
IP: Payload length = 32
IP: NextHeader= 58
IP: Hop limit= 255
IP: Source address = 2001:620:618:197:1:80b2:97c0:1
IP: Destination address = ff02::1:ff01:1
IP:
ICMPv6: ----- ICMPv6 Header -----
ICMPv6:
ICMPv6: Type = 135
ICMPv6: Code=0
ICMPv6: Checksum = 0xb199 [correct]
ICMPv6: Reserved = 00000000
ICMPv6: Target Address=2001:620:618:197:1:80b2:9701:1
ICMPv6:

ETHER: ----- Ether Header -----
ETHER:
ETHER: Packet 2 arrived at 11:55:22.306
ETHER: Packet size = 86 bytes
ETHER: Destination = 3c:07:54:3e:ab:f2
ETHER: Source = 00:08:e3:ff:fc:50
ETHER: Ethertype = 0x86dd
ETHER:
IP: ----- IP Header -----
IP:
IP: Version = 6
IP: Traffic class = 0x000000e0
IP: .... 1110 00.. .... .... .... .... = Class Sector 7 Differentiated Service Field
IP: .... .... ..0. .... .... .... .... = No ECN-Capable Transport (ECT)
IP: .... .... ...0 .... .... .... .... =No ECN-CE
IP: .... .... 0000 0000 0000 0000 0000 = Flowlabel: 0x00000000
IP: Payload length = 32
IP: NextHeader=58 (ICMPv6)
IP: Hop limit= 255
IP: Source address = 2001:620:618:197:1:80b2:9701:1
IP: Destination address = 2001:620:618:197:1:80b2:97c0:1
IP:
ICMPv6: ----- ICMPv6 Header -----
ICMPv6:
```



```
ICMPv6: Type = 136
ICMPv6: Code=0
ICMPv6: Checksum = 0xe3f8 [correct]
ICMPv6: Flags=0xe0000000
ICMPv6: 1... .. = Router: Set
ICMPv6: .1.. .. = Solicited: Set
ICMPv6: ..1. .. = Override: Set
ICMPv6: ...0 0000 0000 0000 0000 0000 0000 0000 = Reserved: 0
ICMPv6: Target Address=2001:620:618:197:1:80b2:9701:1
ICMPv6:
```

Q47/ What are these packets used for in the exchange? What do Types 135 and 136 in the ICMPv6 header represent?

Solution. *Packet 1 is a multicast packet sent by a host seeking the MAC address of a neighbor. Packet 2 is a unicast response packet from the target neighbor that contains the MAC address requested in Packet 1. Type 135 indicates that the packet is a neighbor solicitation packet and Type 136 indicates the packet is a neighbor advertisement packet.*

Q48/ By observing each layer header separately, how do we know that the packet exchange uses IPv6 in each of them?

Solution. *Ethertype = 0x86dd. IP version = 6. ICMPv6*

Shortly afterwards, we observe the following exchange.

```
ETHER: ----- Ether Header -----
ETHER:
ETHER: Packet 3 arrived at 11:55:23.186
ETHER: Packet size = 91 bytes
ETHER: Destination = 00:08:e3:ff:fc:50
ETHER: Source      = 3c:07:54:3e:ab:f2
ETHER: Ethertype = 0x86dd
ETHER:
IP: ----- IP Header -----
IP:
IP: Version = 6
IP: Traffic class = 0x00000000
IP:      .... 0000 00.. .... .... .... = Default Differentiated Service Field
IP:      ....      ..0. .... .... .... = No ECN-Capable Transport (ECT)
IP:      ....      ...0 .... .... .... = No ECN-CE
IP: .... .... .... 0000 0000 0000 0000 0000 = Flowlabel: 0x00000000
IP: payload length = 37
IP: Next header = 17 (UDP)
IP: Hop limit = 255
IP: Source address = 2001:620:618:197:1:80b2:97c0:1
IP: Destination address = 2001:620:618:1a6:1:80b2:a66a:1
IP:
UDP: ----- UDP Header -----
UDP:
UDP: Source port = 63736
UDP: Destination port = 53 (DNS)
UDP: Length = 37
UDP: Checksum =0x4cd7
UDP:
DNS: ----- DNS Header -----
DNS:
DNS: Transaction ID=0x2f34
DNS: Flags= 0x100
DNS:  0...      .... .... = Message is a query
DNS:  .000 0...      .... = Standard query
DNS:      .... ..0. .... = Message not truncated
DNS:      .... ...1 .... = Do query recursively
DNS:      ....      .0.. .... = reserved
DNS:      ....      ...0 .... = Non-authenticated data unacceptable
DNS: Queries
DNS:      ..... Name=www.ethz.ch
DNS:      ..... Record Type= AAAA (IPv6 address)
DNS:      ..... Class = IN (0x0001)
DNS:
```

```

ETHER: ----- Ether Header -----
ETHER:
ETHER: Packet 4 arrived at 11:55:23:191
ETHER: Packet size = 331 bytes
ETHER: Destination = 3c:07:54:3e:ab:f2
ETHER: Source       = 00:08:e3:ff:fc:50
ETHER: Ethertype = 086dd
ETHER:
IP: ----- IP Header -----
IP:
IP: Version = 6
IP: Traffic class = 0x00000000
IP:      .... 0000 00.. .... .... .... .... = Default Differentiated Service Field
IP:      ....      ..0. .... .... .... .... = No ECN-Capable Transport (ECT)
IP:      ....      ...0 .... .... .... .... = No ECN-CE
IP: .... .... .... 0000 0000 0000 0000 0000 = Flowlabel: 0x00000000
IP: Payload length = 277
IP: Next header = 17 (UDP)
IP: Hop limit = 60
IP: Source address = 2001:620:618:1a6:1:80b2:a66a:1
IP: Destination address = 2001:620:618:197:1:80b2:97c0:1
IP:
UDP: ----- UDP Header -----
UDP:
UDP: Source port = 53
UDP: Destination port = 63736
UDP: Length = 277
UDP: Checksum = 0x4a49
UDP:
DNS: ----- DNS Header -----
DNS:
DNS: Flags = 0x8180
DNS:  1...      .... .... ....      = Message is a response
DNS:  .000 0...      .... ....      = Standard query
DNS:      .... .0.. .... ....      = Server not authoritative
DNS:      .... ..0. .... ....      = Message not truncated
DNS:      .... ...1 .... ....      = Do query recursively
DNS:      .... .... 1... ....      = Recursion available
DNS:      .... .... .0.. ....      = reserved
DNS:      .... .... ..0 ....      = Answer/authority portion not authenticated
DNS:      .... .... ...0 ....      = Non-authenticated data unacceptable
DNS:      .... .... .... 0000      = No error
DNS: Queries
DNS:      ..... Name=www.ethz.ch
DNS:      ..... Record Type= AAAA (IPv6 address)
DNS:      ..... Class = IN
DNS: Answers
DNS:      ..... Name=www.ethz.ch
DNS:      ..... Record Type= AAAA
DNS:      ..... Class = IN (0x0001)
DNS:      ..... TTL=1 hour
DNS:      ..... Data length= 16
DNS:      ..... Address = 2001:67c:10ec:4380::216
DNS: Non-authoritative nameservers
DNS:      ..... ""
DNS: Additional records
DNS:      ..... ""
DNS:

```

Q49/ What does this exchange correspond to? Explain both packets.

Solution. First packet: `lrcpc3` issues an UDP packet searching for the IP address of `www.v6.facebook.com`. This packet is sent to the DNS server.
Second packet: DNS server answers with the requested IP address.

Q50/ Mention two ways that can be used to recognize a DNS packet.

Solution. UDP port 53 is reserved for Name Server in DNS.
The DNS header in the body of the UDP packet.

Q51/ What does "non-authoritative nameservers" mean?

Solution. It means that the nameserver which replied to the DNS query got the information second-hand from another nameserver.

For each domain name, there is one or a few authoritative nameservers. For example, Google have a few nameservers that can send authoritative replies to DNS queries concerning Google domain names. If you query the EPFL nameserver for a Google domain, then you will receive a non-authoritative reply since the EPFL nameserver will request this information from the Google nameserver (or have it stored in cache).

Q52/ How would you get the IPv6 address of `www.ethz.ch` using `nslookup`? Try it on your VM and report your findings.

Solution. `nslookup -query=AAAA www.ethz.ch`